



GRA4150

AI - Technologies and Applications

Lecture 7

February 21, 2023

- ▶ Logistic regression;
 - ▶ The sigmoid activation function;
 - ▶ The cross-entropy loss function;
 - ▶ Example on the notebook;

- ▶ Multi-class classification:
 - ▶ One-versus-All;
 - ▶ Softmax regression;
 - ▶ Example on the notebook.

RECAP



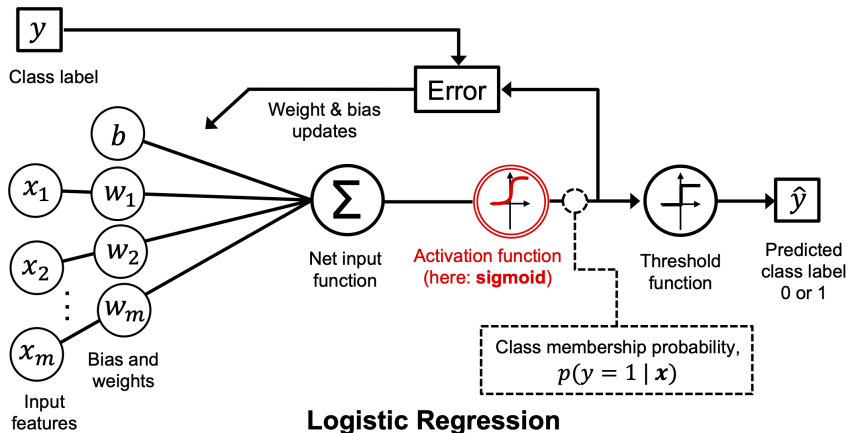
- ▶ We considered **supervised ML** models for **binary classification**;

- ▶ We considered **supervised ML** models for **binary classification**;
- ▶ We started with the PERCEPTRON MODEL:
 - ▶ Not really a motivation for the updating rule (it is sort of "given");
 - ▶ The output \hat{y} is obtained by applying a **threshold function** directly to the net input;
 - ▶ The updating step is computed from the final output \hat{y} and the observation y ;
 - ▶ Notice that the convergence of the method can be proved mathematically.

- ▶ We considered **supervised ML** models for **binary classification**;
- ▶ We started with the PERCEPTRON MODEL:
 - ▶ Not really a motivation for the updating rule (it is sort of "given");
 - ▶ The output \hat{y} is obtained by applying a **threshold function** directly to the net input;
 - ▶ The updating step is computed from the final output \hat{y} and the observation y ;
 - ▶ Notice that the convergence of the method can be proved mathematically.
- ▶ The ADALINE METHOD:
 - ▶ It is an extension of the previous model;
 - ▶ Here we **add an intermediate step**: an **activation function** is applied to the net input to get the intermediate output \tilde{y} ;
 - ▶ The updating step is computed from the intermediate output \tilde{y} and observation y ;
 - ▶ In this case, the activation function is in fact a **linear activation function**;
 - ▶ The benefit of this approach is that the linear activation function is **differentiable**;
 - ▶ We can then define an **objective function** (the MSE loss function) and use a **gradient method** for the optimization;
 - ▶ Only in the very last step, the intermediate output \tilde{y} is converted in final output \hat{y} via a threshold function.

- ▶ Logistic regression is also a model for binary classification;
- ▶ It is one of the most widely used algorithms for classification in industry;
- ▶ It is a **probabilistic model**;
- ▶ It easily generalizes to multi-class classification.

THE LOGISTIC REGRESSION ALGORITHM




What is this similar to? What is/are the difference(s)?

THE SIGMOID FUNCTION

We define the **logistic sigmoid function**, or simply **sigmoid function** by:

$$\sigma(z) := \frac{1}{1 + e^{-z}}$$

where $z = \mathbf{w}^\top \mathbf{x} + b$ is the net input.



Notice that:

- ▶ For z approaching $+\infty$, $\sigma(z)$ approaches 1;
- ▶ For z approaching $-\infty$, $\sigma(z)$ approaches 0;
- ▶ Hence $\sigma : \mathbb{R} \rightarrow [0, 1]$ with intercept $\sigma(0) = 0.5$.

WHY THE SIGMOID FUNCTION?

Let $p \in [0, 1]$ the probability of a "positive event" (here "positive" does not mean "good", but refers to the event that we want to predict).

Example

We consider p as the probability that a patient has a certain disease given certain symptoms. Then we think of the positive event as class label $y = 1$ and the symptoms as feature \mathbf{x} , and $p := p(y = 1|\mathbf{x})$ is the probability that a particular example belongs to the class 1 given \mathbf{x} .

The output of the sigmoid function is interpreted as the probability of a particular example belonging to class 1:

$$\sigma(z) = p(y = 1|\mathbf{x}; \mathbf{w}, b)$$

given its features \mathbf{x} and parameterized by the weights and bias, \mathbf{w} and b .

Example

If $\sigma(z) = 0.8$ for a particular flower example, it means that the chance that this example is an Iris-versicolor flower is 80%. Therefore, the probability that this flower is an Iris-setosa flower can be calculated as $p(y = 0|\mathbf{x}; \mathbf{w}, b) = 1 - p(y = 1|\mathbf{x}; \mathbf{w}, b) = 0.2$, or 20%.

THE (same) THRESHOLD FUNCTION (as Adaline)

The predicted probability can then simply be converted into a binary outcome via a threshold function:

$$\hat{y} = \begin{cases} 1 & \text{if } \sigma(z) \geq 0.5 \\ 0 & \text{if } \sigma(z) < 0.5 \end{cases}$$

REMARK

There are many applications where we are interested in both the predicted class labels and in **the estimation of the class-membership probability** (the output of the sigmoid function prior to applying the threshold function).

Examples:

- ▶ Weather forecasting: we are interested not only to predict whether it will rain on a particular day, but also to report the chance of rain.
- ▶ Medicine: logistic regression can be used to predict the chance that a patient has a particular disease given certain symptoms.

WHAT ABOUT THE LEARNING PROCESS?

- ▶ We need to define an objective function that we seek to minimize;
- ▶ The idea is that we want to **maximize the likelihood**, namely the probability of classifying the dataset correctly, given some parameters \mathbf{w} and b ;
- ▶ This is equivalent to minimize the **CROSS-ENTROPY loss function**:

$$L(\mathbf{w}, b) := \sum_{i=1}^n \left[-y^{(i)} \log \left(\sigma(z^{(i)}) \right) - (1 - y^{(i)}) \log \left(1 - \sigma(z^{(i)}) \right) \right].$$

Example: $n = 1$


$$L(\mathbf{w}, b) = -y^{(1)} \log \left(\sigma(z^{(1)}) \right) - (1 - y^{(1)}) \log \left(1 - \sigma(z^{(1)}) \right)$$

We notice that:

$$L(\mathbf{w}, b) = \begin{cases} -\log \left(\sigma(z^{(1)}) \right) & \text{if } y^{(1)} = 1 \\ -\log \left(1 - \sigma(z^{(1)}) \right) & \text{if } y^{(1)} = 0 \end{cases}.$$

THE CROSS-ENTROPY FUNCTION EXPLAINED

$$L(\mathbf{w}, b) = \begin{cases} -\log(\sigma(z^{(1)})) & \text{if } y^{(1)} = 1 \\ -\log(1 - \sigma(z^{(1)})) & \text{if } y^{(1)} = 0 \end{cases}$$



- ▶ Blue continuous line: the loss approaches 0 if we correctly predict that an example belongs to class 1;
- ▶ Orange dashed line: the loss also approaches 0 if we correctly predict $y = 0$;
- ▶ If the prediction is wrong, the loss goes toward infinity;
- ▶ **We penalize wrong predictions with an increasingly larger loss!**

The logistic regression algorithm can be obtained starting from the Adaline implementation:

1. By substituting the MSE loss function with the cross-entropy loss function;
2. By changing the linear activation function with the sigmoid function;

► What about the gradient for the optimization step?

The logistic regression algorithm can be obtained starting from the Adaline implementation:

1. By substituting the MSE loss function with the cross-entropy loss function;
 2. By changing the linear activation function with the sigmoid function;
-
- ▶ What about the gradient for the optimization step?
 - ▶ Using calculus, we can show (optional exercise at home) that **the parameter updates via gradient descent are equal to the ones for Adaline**, namely the gradient step is unchanged!

THE GRADIENT STEP (unchanged)

- By **gradient descent**, we update the model parameters by taking a step in the opposite direction of the gradient of the loss function:

$$\mathbf{w} = \mathbf{w} + \Delta \mathbf{w}$$

$$\text{with } \Delta \mathbf{w} = -\eta \nabla_{\mathbf{w}} L(\mathbf{w}, b)$$

$$b = b + \Delta b$$

$$\text{with } \Delta b = -\eta \nabla_b L(\mathbf{w}, b)$$

- $\nabla_{\mathbf{w}} L(\mathbf{w}, b)$ is the vector whose components are the partial derivatives of the loss function with respect to each weight w_j :

$$\nabla_{\mathbf{w}} L(\mathbf{w}, b) = \left(\frac{\partial L}{\partial w_1} \quad \frac{\partial L}{\partial w_2} \quad \cdots \quad \frac{\partial L}{\partial w_m} \right)^\top \text{ with } \frac{\partial L}{\partial w_j} = -\frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - \sigma(z^{(i)}) \right) x_j^{(i)};$$

- $\nabla_b L(\mathbf{w}, b)$ corresponds to the partial derivative of the loss function with respect to the bias b :

$$\nabla_b L(\mathbf{w}, b) = \frac{\partial L}{\partial b} = -\frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - \sigma(z^{(i)}) \right).$$

- ▶ Logistic regression works for binary classification;
- ▶ We consider a different activation function σ , the sigmoid function;
- ▶ We consider a different loss function L , the cross-entropy function;
- ▶ It has the additional property of providing the class-membership likelihood/probability: this is the output of the sigmoid activation function.

- ▶ Logistic regression works for binary classification;
- ▶ We consider a different activation function σ , the sigmoid function;
- ▶ We consider a different loss function L , the cross-entropy function;
- ▶ It has the additional property of providing the class-membership likelihood/probability: this is the output of the sigmoid activation function.

What about multi-class classification?

- ▶ Logistic regression works for binary classification;
- ▶ We consider a different activation function σ , the sigmoid function;
- ▶ We consider a different loss function L , the cross-entropy function;
- ▶ It has the additional property of providing the class-membership likelihood/probability: this is the output of the sigmoid activation function.


What about multi-class classification?

We can consider two different approaches:

- ▶ The One-vs-All method;
- ▶ The multinomial logistic regression.

FIRST APPROACH: THE OvA METHOD FOR MULTI-CLASS CLASSIFICATION

- ▶ The OvA (= One-versus-All) method is a technique that allows to extend **probabilistic binary classifiers** to multi-class problems;
- ▶ The idea is to **train one classifier per class**, where the particular class is treated as the positive class ($y = 1$) and the examples from all other classes are considered negative classes ($y = 0$);
- ▶ If there are N classes, then **we train N different classifiers**;
- ▶ We choose the class label with the highest confidence.




We have to generate the same number of classifiers as the number of class labels:

- ▶ Classifier 1: [Green] vs [Red, Blue], answers the question "Is it green?"
- ▶ Classifier 2: [Blue] vs [Green, Red], answers the question "Is it blue?"
- ▶ Classifier 3: [Red] vs [Blue, Green], answers the question "Is it red?"

Important: to train these three classifiers, we need to create three training datasets.

After training, we pass a test data to the model:

- ▶ Suppose that from the first classifier we get a probability score of $p = 0.8$;
(it means probability of being Green 0.8 and probability of not being Green 0.2)
- ▶ From the second classifier we get a probability score of $p = 0.4$;
(it means probability of being Blue 0.4 and probability of not being Blue 0.6)
- ▶ From the third classifier we get a probability score of $p = 0.35$;
(it means probability of being Red 0.35 and probability of not being Red 0.65)
- ▶ We can say that the test input belongs to the Green class.



Every classifier returns a class probability

$$\sigma(z_h) = p_h \quad \text{with} \quad z_h := \mathbf{w}_h^\top \mathbf{x} + b_h \quad \text{for} \quad h = 1, \dots, N.$$

Here each \mathbf{w}_h is a vector with components $\mathbf{w}_h^\top = (w_{h,1}, w_{h,2}, \dots, w_{h,m})$.

OvA IN MATRIX FORM

Instead of writing for $h = 1, \dots, N$,

$$\sigma(z_h) = p_h \text{ with } z_h := \mathbf{w}_h^\top \mathbf{x} + b_h,$$

we can use the compact form


$$\sigma(\mathbf{z}) = \mathbf{p} \text{ with } \mathbf{z} := W\mathbf{x} + \mathbf{b}$$

where

- ▶ $\mathbf{z} \in \mathbb{R}^N$ is the vector with components z_h ;
- ▶ $\mathbf{p} \in \mathbb{R}^N$ is the vector with components p_h ;
- ▶ $\mathbf{b} \in \mathbb{R}^N$ is the vector with components b_h ;
- ▶ $W \in \mathbb{R}^{N \times m}$ is the matrix whose rows are the vectors \mathbf{w}_h^\top ;
- ▶ ...and N is the number of classes.
- ▶ Notice that with this notation we mean that σ acts component-wise, namely $\sigma(\mathbf{z}) = (\sigma(z_1), \dots, \sigma(z_N))^\top$.

SECOND APPROACH: SOFTMAX REGRESSION

Multinomial logistic regression



THE SOFTMAX ACTIVATION FUNCTION

The softmax function $\sigma : \mathbb{R}^N \rightarrow [0, 1]^N$ is defined as

$$\sigma(\mathbf{z})_h = \frac{e^{z_h}}{\sum_{\ell=1}^N e^{z_\ell}},$$

where $\sigma(\mathbf{z})_h$ is the h -th component of $\sigma(\mathbf{z}) = \sigma(W\mathbf{x} + \mathbf{b})$, for $h = 1, \dots, N$.

In words:

- ▶ It applies the exponential function to each element z_h of the input vector \mathbf{z} ;
- ▶ It normalizes these values by dividing by the sum of all the exponentials;
- ▶ The normalization ensures that the sum of the components of $\sigma(\mathbf{z})$ is 1;
- ▶ This implies that all the components of $\sigma(\mathbf{z})$ take values between 0 and 1;
- ▶ They indeed represent probabilities:

$$\sigma(\mathbf{z})_h = p_h = p(y = h | \mathbf{x}; W, \mathbf{b}).$$

- ▶ Softmax returns independent probabilities;
- ▶ Softmax returns in fact a global distribution where the sum of all the probabilities equal to 1;
- ▶ This is not the case for One-vs-All;
- ▶ The big difference is that **Softmax assumes that each example is a member of exactly one class**. Some examples, however, can simultaneously be a member of multiple classes (think of an image with both a pear and an apple).
- ▶ One-vs-All is reasonable when the total number of classes is small, but becomes inefficient as the number of classes rises.

Notice also that:

- ▶ We have now covered the case of multi-class classification;
- ▶ With the logistic regression, we have added "non-linearity" to the model because of the sigmoid function;
- ▶ This allows us to also threat non-linearly separable classes.

- ▶ We implemented the Adaline algorithm to perform **binary classification**;
- ▶ We used the **gradient descent** optimization algorithm to learn the weight and bias coefficients of the model;
- ▶ The function to optimize was the MSE (**Mean Squared Error**) loss function;
- ▶ The weights were updated by taking a step in the opposite direction of the gradient, with the learning rate η as a scale factor;
- ▶ We updated the weights simultaneously after each epoch;
- ▶ For the final prediction, we introduced a **threshold function**;
- ▶ By changing loss function and activation function, we easily obtain the **logistic regression** model: their structure is essentially the same!!!

Both Adaline and logistic regression can be referred to as **single-layer neural network**.

- ▶ We obtained multi-class classification by rewriting the problem in **matrix form**.

1. Lecture 8: *On robustness, hallucinations and safety of AI algorithms* with V. Antun (UiO);
2. Lecture 9:
 - ▶ Discussion Exercise 2;
 - ▶ Discussion Exercise 3;
3. Lecture 10:
 - ▶ Neural Networks: backpropagation and convergence;
 - ▶ Guest lecture with L. Pararasasingam (Statkraft);
4. Lecture 11:
 - ▶ Neural Networks: Stochastic Gradient Descent;
 - ▶ Discussion Exercise 4.
5. Lecture 12:
 - ▶ TBD.